

▼ Part 1 : Load Dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force_re

```
!pip install transformers
!pip install pytorch_wrapper
!pip install seaborn
!nvidia-smi
```

```
from google.colab import drive
drive.mount('/content/drive')
!unrar x -Y /content/drive/MyDrive/dataset.rar dataset.csv
!grep -v is_fake dataset.csv > clean_dataset.csv
!git clone https://github.com/nlpaueb/greek-bert.git
```

```
Requirement already satisfied: transformers in /usr/local/lib/python3.7/dist-packages (4.11.3)
Requirement already satisfied: pyyaml<=5.1 in /usr/local/lib/python3.7/dist-packages (from transformers) (6.0)
Requirement already satisfied: tqdm<=4.27 in /usr/local/lib/python3.7/dist-packages (from transformers) (4.62.3)
Requirement already satisfied: sacremoses in /usr/local/lib/python3.7/dist-packages (from transformers) (0.0.46)
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.7/dist-packages (from transformers) (2020.1.8)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.7/dist-packages (from transformers) (21.0)
Requirement already satisfied: importlib-metadata in /usr/local/lib/python3.7/dist-packages (from transformers) (4.2.0)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.7/dist-packages (from transformers) (1.19.5)
Requirement already satisfied: tokenizers<0.11,>=0.10.1 in /usr/local/lib/python3.7/dist-packages (from transformers) (0.10.2)
Requirement already satisfied: filelock in /usr/local/lib/python3.7/dist-packages (from transformers) (3.3.0)
Requirement already satisfied: huggingface-hub>=0.0.17 in /usr/local/lib/python3.7/dist-packages (from transformers) (0.5.0)
Requirement already satisfied: requests in /usr/local/lib/python3.7/dist-packages (from transformers) (2.23.0)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from huggingface-hub>=0.0.17) (3.7.4)
Requirement already satisfied: pyparsing>=2.0.2 in /usr/local/lib/python3.7/dist-packages (from packaging>=20.0->transformers) (2.4.7)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from importlib-metadata->transformers) (3.4.0)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests->transformers) (2020.12.5)
```

```
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests->transformers)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests->transformers)
Requirement already satisfied: urllib3!=1.25.0,!>=1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests->transformers)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from sacremoses->transformers) (1.15.0)
Requirement already satisfied: click in /usr/local/lib/python3.7/dist-packages (from sacremoses->transformers) (7.0)
Requirement already satisfied: joblib in /usr/local/lib/python3.7/dist-packages (from sacremoses->transformers) (1.0.1)
Requirement already satisfied: pytorch_wrapper in /usr/local/lib/python3.7/dist-packages (1.1.0)
Requirement already satisfied: six<2,>=1.11 in /usr/local/lib/python3.7/dist-packages (from pytorch_wrapper) (1.15.0)
Requirement already satisfied: numpy<2,>=1.14 in /usr/local/lib/python3.7/dist-packages (from pytorch_wrapper) (1.19.5)
Requirement already satisfied: scikit-learn<1,>=0.19 in /usr/local/lib/python3.7/dist-packages (from pytorch_wrapper) (0.22.2)
Requirement already satisfied: tqdm<5,>=4.36.1 in /usr/local/lib/python3.7/dist-packages (from pytorch_wrapper) (4.62.3)
Requirement already satisfied: torch<2,>=1 in /usr/local/lib/python3.7/dist-packages (from pytorch_wrapper) (1.9.0)
Requirement already satisfied: hyperopt<0.2,>=0.1 in /usr/local/lib/python3.7/dist-packages (from pytorch_wrapper) (0.1.2)
Requirement already satisfied: pymongo in /usr/local/lib/python3.7/dist-packages (from hyperopt<0.2,>=0.1->pytorch_wrapper) (3.10.0)
Requirement already satisfied: future in /usr/local/lib/python3.7/dist-packages (from hyperopt<0.2,>=0.1->pytorch_wrapper) (0.18.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages (from hyperopt<0.2,>=0.1->pytorch_wrapper) (1.5.4)
Requirement already satisfied: networkx in /usr/local/lib/python3.7/dist-packages (from hyperopt<0.2,>=0.1->pytorch_wrapper) (2.6.3)
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.7/dist-packages (from scikit-learn<1,>=0.19->pytorch_wrapper) (1.0.1)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch<2,>=1->pytorch_wrapper) (3.7.4)
Requirement already satisfied: seaborn in /usr/local/lib/python3.7/dist-packages (0.11.2)
Requirement already satisfied: matplotlib>=2.2 in /usr/local/lib/python3.7/dist-packages (from seaborn) (3.2.2)
Requirement already satisfied: numpy>=1.15 in /usr/local/lib/python3.7/dist-packages (from seaborn) (1.19.5)
Requirement already satisfied: scipy>=1.0 in /usr/local/lib/python3.7/dist-packages (from seaborn) (1.4.1)
Requirement already satisfied: pandas>=0.23 in /usr/local/lib/python3.7/dist-packages (from seaborn) (1.1.5)
Requirement already satisfied: python-dateutil>=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.2->seaborn) (2.8.1)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.2->seaborn) (1.3.1)
Requirement already satisfied: pyparsing!=2.0.4,!>=2.1.2,!>=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.2->seaborn) (2.4.7)
Requirement already satisfied: cycycler>=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib>=2.2->seaborn) (0.10.0)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from cycycler>=0.10->matplotlib>=2.2->seaborn) (1.15.0)
Requirement already satisfied: pytz>=2017.2 in /usr/local/lib/python3.7/dist-packages (from pandas>=0.23->seaborn) (2019.3)
Wed Oct 27 09:29:07 2021
```

NVIDIA-SMI 470.74				Driver Version: 460.32.03		CUDA Version: 11.2	
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC	
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	
						MIG M.	
0	Tesla K80	Off	00000000:00:04.0	Off			0
N/A	72C	P0	75W / 149W	9740MiB / 11441MiB	0%	Default	N/A

```
import datetime
import os
import random
import time
import unicodedata

import numpy as np
import pandas as pd
import torch
import matplotlib
import matplotlib.pyplot as plt
import transformers
from transformers import AutoTokenizer, AutoModel, get_linear_schedule_with_warmup
import pytorch_wrapper.functional as pwF
import sklearn
import seaborn
```

```
plt.rcParams['figure.figsize'] = [20, 5]
```

```
print("Seaborn version: %s" % (seaborn.__version__))
print("Sklearn version: %s" % (sklearn.__version__))
print("Transformers version: %s" % (transformers.__version__))
print("PyTorch version: %s" % (torch.__version__))
print("Matplotlib version: %s" % (matplotlib.__version__))
print("NumPy version: %s" % (np.__version__))
print("Pandas version: %s" % (pd.__version__))
```


```
Seaborn version: 0.11.2
Sklearn version: 0.22.2.post1
Transformers version: 4.11.3
PyTorch version: 1.9.0+cu111
Matplotlib version: 3.2.2
NumPy version: 1.19.5
Pandas version: 1.1.5
```

```
df = pd.read_csv("clean_dataset.csv")
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 34976 entries, 0 to 34975
Data columns (total 5 columns):
#   Column   Non-Null Count  Dtype
---  -
0   id        34976 non-null  int64
1   title     34970 non-null  object
2   author    0 non-null      float64
3   text      34611 non-null  object
4   label     34976 non-null  int64
dtypes: float64(1), int64(2), object(2)
memory usage: 1.3+ MB
```

```
#plot labels
labels = df.label
plt.hist(labels)
plt.xlabel("label")
plt.ylabel("number of samples")
plt.title("Labels Distribution")
plt.xticks(np.arange(len(np.unique(labels))))
```

```
([<matplotlib.axis.XTick at 0x7fd14a4229d0>,
 <matplotlib.axis.XTick at 0x7fd14a422990>],
 <a list of 2 Text major ticklabel objects>)
```



```
# plot article lengths and discarded data length
lengths = pd.DataFrame([len(x.split()) for x in df.title + df.text if isinstance(x, str)])
lost_data = pd.DataFrame([max(len(x.split())-512,0) for x in df.title + df.text if isinstance(x, str)])
lost_data_ratios = pd.DataFrame([max(len(x.split())-512,0)/float(len(x.split())) for x in df.title + df.text if isinstance(x, str)])

print("Article lengths description: %s\n" % (lengths.describe()))
lengths.plot(xlabel="article id", ylabel="word count", title="Article Length in Words", legend=None)
print("Article discarded data description: %s\n" % (lost_data.describe()))
lost_data.plot(xlabel="article id", ylabel="words removed", title="Number of Deleted Words per Article", legend=None)
print("Article discarded data ratios description: %s\n" % (lost_data_ratios.describe()))
lost_data_ratios.plot(xlabel="article id", ylabel="words removed ratio", title="Ratio of Deleted Words per Article", legend=None)

s1 = lost_data.sum()
s2 = lengths.sum()
ratio = float(s1) / float(s2)
s3 = (lengths <= 512).astype(int).sum(axis=0)
ratio2 = float(s3) / float(len(lengths))
print("Articles exceeding 512 word limit: %d out of %d (%.2lf%%)" % (s3, len(lengths), ratio2*100))
print("Total data loss: %d out of %d words (%.2lf%% total, %.2lf%% average)" % (s1, s2, ratio*100, lost_data_ratios.mean()))
```

Article lengths description: 0

count 34605.00
mean 347.27
std 551.12
min 5.00
25% 128.00
50% 227.00
75% 385.00
max 42891.00

Article discarded data description: 0


count 34605.00
mean 89.87
std 482.38
min 0.00
25% 0.00
50% 0.00
75% 0.00
max 42379.00



Article discarded data ratios description: 0

count 34605.00
mean 0.06
std 0.17
min 0.00
25% 0.00
50% 0.00
75% 0.00
max 0.99

Articles exceeding 512 word limit: 28942 out of 34605 (83.64%)

Total data loss: 3109972 out of 12017277 words (25.88% total, 6.04% average)





▼ Part 2: Tokenization and Input Formatting

```
start = time.time()

def strip_accents_and_lowercase(s):
    return ''.join(c for c in unicodedata.normalize('NFD', s)
                    if unicodedata.category(c) != 'Mn').lower()
#####
class BinaryGreekBERTModel(torch.nn.Module):

    def __init__(self, bert_model, dp):
        super(BinaryGreekBERTModel, self).__init__()
        self._bert_model = bert_model
        self._dp = torch.nn.Dropout(dp)
        self._output_linear = torch.nn.Linear(768, 1)

    def forward(self, text, text_len):
        attention_mask = pwF.create_mask_from_length(text_len, text.shape[1])
        return self._output_linear(self._dp(self._bert_model(text, attention_mask=attention_mask)[0][:, 0, :]))
#####

#from transformers import BertTokenizer
#tokenizer = BertTokenizer.from_pretrained('bert-base-uncased', do_lower_case=True)
tokenizer = AutoTokenizer.from_pretrained("nlpueb/bert-base-greek-uncased-v1")
greekbert_model = AutoModel.from_pretrained('nlpueb/bert-base-greek-uncased-v1')
binary_greekbert_model = BinaryGreekBERTModel(greekbert_model, dp=0)

working_set_df = df.head(2000)

text_data = working_set_df.title + working_set_df.text
text_data.fillna('', inplace=True)
```



```

input_ids = []
attention_masks = []
for item in text_data:
    unaccented_string = strip_accents_and_lowercase(item)
    encoded_dict = tokenizer.encode_plus(unaccented_string,
                                        add_special_tokens=True,
                                        max_length=512,
                                        truncation=True,
                                        padding='max_length',
                                        #pad_to_max_length=True,
                                        return_attention_mask=True,
                                        return_tensors="pt")

    input_ids.append(encoded_dict["input_ids"])
    attention_masks.append(encoded_dict["attention_mask"])

input_ids = torch.cat(input_ids, dim=0)
attention_masks = torch.cat(attention_masks, dim=0)
labels = torch.tensor(working_set_df.label.values)


```

```

print("Time elapsed: %.1lf seconds." % (time.time() - start))

```

Some weights of the model checkpoint at nlpauieb/bert-base-greek-uncased-v1 were not used when initializing BertModel
- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactl
Time elapsed: 9.6 seconds.



```

# split training and validation sets
from torch.utils.data import TensorDataset, random_split, DataLoader, RandomSampler, SequentialSampler

dataset = TensorDataset(input_ids, attention_masks, labels)
train_size = int(0.8 * len(dataset))
validation_size = int(0.1 * len(dataset))
test_size = int(0.1 * len(dataset))
train_dataset, validation_dataset, test_dataset = random_split(dataset, [train_size, validation_size, test_size])
print("Training size: %s, Validation size: %s Test size: %s" % (train_size, validation_size, test_size))

```

```
batch_size = 8
train_dataloader = DataLoader(train_dataset, sampler=RandomSampler(train_dataset), batch_size=batch_size)
validation_dataloader = DataLoader(validation_dataset, sampler=SequentialSampler(validation_dataset), batch_size=batch_si
test_dataloader = DataLoader(test_dataset, sampler=SequentialSampler(test_dataset), batch_size=batch_size)
```

Training size: 1600, Validation size: 200 Test size: 200

▼ Part 3: Training

```
start = time.time()

from transformers import BertForSequenceClassification, AdamW, BertConfig

model = BertForSequenceClassification.from_pretrained(
    "nlpau/bert-base-greek-uncased-v1",
    num_labels = 2,
    output_attentions = False,
    output_hidden_states = False,
)

device = None
if torch.cuda.is_available():
    device = torch.device("cuda:0")
    model.cuda()
else:
    device = torch.device("cpu")

#####
# Batch sizes: 16, 32
# Learning rate (Adam): 5e-5, 3e-5, 2e-5
# Epochs: 2, 3, 4
#####

optimizer = AdamW(model.parameters(), lr=2e-5, eps=1e-8)
```

```

epochs = 4
total_steps = len(train_dataloader) * epochs
scheduler = get_linear_schedule_with_warmup(optimizer,
                                             num_warmup_steps = 0, # Default value in run_glue.py
                                             num_training_steps = total_steps)

```

```

print("Time elapsed: %.2lf" % (time.time() - start))
!nvidia-smi

```

Some weights of the model checkpoint at nlpaueb/bert-base-greek-uncased-v1 were not used when initializing BertForSequenceClassification from the checkpoint of a model trained on a different task. This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on a different task. This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you expect to be initialized from the checkpoint of a model trained on this task. Some weights of BertForSequenceClassification were not initialized from the model checkpoint at nlpaueb/bert-base-greek-uncased-v1 and are newly initialized from random normal distribution. You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Time elapsed: 1.81
Wed Oct 27 09:00:36 2021

+-----+-----+-----+-----+-----+-----+									
NVIDIA-SMI 470.74			Driver Version: 460.32.03			CUDA Version: 11.2			
+-----+-----+-----+-----+-----+-----+									
GPU	Name		Persistence-M	Bus-Id	Disp.A	Volatile Uncorr. ECC			
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute M.		
							MIG M.		
=====+=====+=====+=====+=====+=====									
0	Tesla K80		Off	00000000:00:04.0	Off		0		
N/A	70C	P0	76W / 149W	9740MiB / 11441MiB		36%	Default		
							N/A		
+-----+-----+-----+-----+-----+-----+									

+-----+-----+-----+-----+-----+-----+							
Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory	
	ID	ID				Usage	
=====+=====+=====+=====+=====+=====							
No running processes found							
+-----+-----+-----+-----+-----+-----+							

▼ Training Loop

```

start = time.time()

def flat_accuracy(preds, labels):
    """Function to calculate the accuracy of our predictions vs labels"""
    pred_flat = np.argmax(preds, axis=1).flatten()
    labels_flat = labels.flatten()
    return np.sum(pred_flat == labels_flat) / len(labels_flat)
#####

seed_val = 21
random.seed(seed_val)
np.random.seed(seed_val)
torch.manual_seed(seed_val)
torch.cuda.manual_seed_all(seed_val)

training_stats = []
total_t0 = time.time()

for epoch_i in range(0, epochs):
    print('Epoch %s out of %s' % (epoch_i + 1, epochs))
    print('Training...')

    t0 = time.time()
    total_train_loss = 0
    model.train()
    for step, batch in enumerate(train_dataloader):
        if step % 40 == 0 and not step == 0:
            print(' Batch %s out of %s, Elapsed %.2lf seconds' % (step, len(train_dataloader), time.time() - t0))

        b_input_ids = batch[0].to(device)
        b_input_mask = batch[1].to(device)
        b_labels = batch[2].to(device)
        model.zero_grad()
        loss, logits = model(
            b_input_ids,
            token_type_ids=None,

```

```

        attention_mask=b_input_mask,
        labels=b_labels,
        return_dict=False
    )

    total_train_loss += loss.item()
    loss.backward()
    torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
    optimizer.step()
    scheduler.step()
    avg_train_loss = total_train_loss / len(train_dataloader)
    training_time = time.time() - t0

    print(" Average training loss: %.2lf" % (avg_train_loss))
    print(" Epoch %s training time %.2lf seconds" % (epoch_i, training_time))
    print("Running Validation...")

    t0 = time.time()
    model.eval()
    total_eval_accuracy = 0
    total_eval_loss = 0
    nb_eval_steps = 0
    for batch in validation_dataloader:
        b_input_ids = batch[0].to(device)
        b_input_mask = batch[1].to(device)
        b_labels = batch[2].to(device)
        with torch.no_grad():
            loss, logits = model(
                b_input_ids,
                token_type_ids=None,
                attention_mask=b_input_mask,
                labels=b_labels,
                return_dict=False
            )

        total_eval_loss += loss.item()
        logits = logits.detach().cpu().numpy()

```

```

        label_ids = b_labels.to('cpu').numpy()
        total_eval_accuracy += flat_accuracy(logits, label_ids)
    avg_val_accuracy = total_eval_accuracy / len(validation_dataloader)
    avg_val_loss = total_eval_loss / len(validation_dataloader)
    validation_time = time.time() - t0

    print("  Accuracy: %.2lf" % (avg_val_accuracy))
    print("  Validation Loss: %2lf" % (avg_val_loss))
    print("  Validation time %.2lf seconds" % (validation_time))

    training_stats.append(
        {
            'epoch': epoch_i + 1,
            'Training Loss': avg_train_loss,
            'Valid. Loss': avg_val_loss,
            'Valid. Accur.': avg_val_accuracy,
            'Training Time': training_time,
            'Validation Time': validation_time
        }
    )
print("Training complete!")
print("Time elapsed: %.2lf seconds" % (time.time() - start))

```

Epoch 1 out of 4

Training...

```

Batch 40 out of 200, Elapsed 62.23 seconds
Batch 80 out of 200, Elapsed 124.20 seconds
Batch 120 out of 200, Elapsed 186.11 seconds
Batch 160 out of 200, Elapsed 248.16 seconds
Average training loss: 0.37
Epoch 0 training time 310.11 seconds

```

Running Validation...

```

Accuracy: 0.91
Validation Loss: 0.308900
Validation time 13.87 seconds

```

Epoch 2 out of 4

Training...

```

Batch 40 out of 200, Elapsed 61.84 seconds
Batch 80 out of 200, Elapsed 123.63 seconds

```

Batch 120 out of 200, Elapsed 185.44 seconds
Batch 160 out of 200, Elapsed 247.28 seconds
Average training loss: 0.15
Epoch 1 training time 309.16 seconds
Running Validation...
Accuracy: 0.91
Validation Loss: 0.349148
Validation time 13.87 seconds
Epoch 3 out of 4
Training...
Batch 40 out of 200, Elapsed 61.68 seconds
Batch 80 out of 200, Elapsed 123.35 seconds
Batch 120 out of 200, Elapsed 185.04 seconds
Batch 160 out of 200, Elapsed 246.73 seconds
Average training loss: 0.06
Epoch 2 training time 308.48 seconds
Running Validation...
Accuracy: 0.92
Validation Loss: 0.479133
Validation time 13.86 seconds
Epoch 4 out of 4
Training...
Batch 40 out of 200, Elapsed 61.71 seconds
Batch 80 out of 200, Elapsed 123.35 seconds
Batch 120 out of 200, Elapsed 184.84 seconds
Batch 160 out of 200, Elapsed 246.53 seconds
Average training loss: 0.02
Epoch 3 training time 308.13 seconds
Running Validation...
Accuracy: 0.91
Validation Loss: 0.500599
Validation time 13.88 seconds
Training complete!
Time elapsed: 1291.37 seconds

!nvidia-smi

Wed Oct 27 09:25:45 2021

```
+-----+
| NVIDIA-SMI 470.74      Driver Version: 460.32.03   CUDA Version: 11.2   |
|-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+
|
```

GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.
						MIG M.
0	Tesla K80	Off	00000000:00:04.0	Off	0	
N/A	74C	P0	76W / 149W	9740MiB / 11441MiB	0%	Default
						N/A

Processes:

GPU	GI	CI	PID	Type	Process name	GPU Memory
	ID	ID				Usage
No running processes found						

```
#####
# View training summary
#####
```

```
pd.set_option('precision', 2)
df_stats = pd.DataFrame(data=training_stats)
df_stats = df_stats.set_index('epoch')
df_stats
```

	Training Loss	Valid. Loss	Valid. Accur.	Training Time	Validation Time
epoch					
1	0.37	0.31	0.91	310.11	13.87
2	0.15	0.35	0.91	309.16	13.87
3	0.06	0.48	0.92	308.48	13.86
4	0.02	0.50	0.91	308.13	13.88

```
# If validation loss is increasing, it means we are overfitting!
```



```
# Plot training and validation loss
import seaborn as sns

# Use plot styling from seaborn.
sns.set(style='darkgrid')

# Increase the plot size and font size.
sns.set(font_scale=1.5)
plt.rcParams["figure.figsize"] = (12,6)

# Plot the learning curve.
plt.plot(df_stats['Training Loss'], 'b-o', label="Training")
plt.plot(df_stats['Valid. Loss'], 'g-o', label="Validation")

# Label the plot.
plt.title("Training & Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.xticks([1, 2, 3, 4])

plt.show()
```



```
#####
```

```
# Save model to disk
```

```
#####
```

```
tstamp = datetime.datetime.now().strftime("%Y-%m-%d_%H%M%S")
```

```
output_dir = './greekbert_model_save/%s/' % (tstamp)
```

```
if not os.path.exists(output_dir):
```

```
    os.makedirs(output_dir)
```

```
print("Saving model to %s" % output_dir)
```

```
model_to_save = model.module if hasattr(model, 'module') else model # Take care of distributed/parallel training
```

```
model_to_save.save_pretrained(output_dir)
```

```
tokenizer.save_pretrained(output_dir)
```

```
torch.save(optimizer.state_dict(), os.path.join(output_dir, 'training_args.bin'))
```

```
    Saving model to ./greekbert_model_save/2021-10-27_092610/
```

```
!ls -lhatrR greekbert_model_save/
```

```
greekbert_model_save/:
```

```
total 16K
```

```
drwxr-xr-x 1 root root 4.0K Oct 27 08:39 ..
```

```
drwxr-xr-x 2 root root 4.0K Oct 27 08:39 2021-10-27_083917
```

```
drwxr-xr-x 4 root root 4.0K Oct 27 09:26 .
```

```
drwxr-xr-x 2 root root 4.0K Oct 27 09:26 2021-10-27_092610
```

```
greekbert_model_save/2021-10-27_083917:
```

```
total 1.3G
```

```
-rw-r--r-- 1 root root 732 Oct 27 08:39 config.json
```

```
-rw-r--r-- 1 root root 526 Oct 27 08:39 tokenizer_config.json
```

```
-rw-r--r-- 1 root root 112 Oct 27 08:39 special_tokens_map.json
```

```
-rw-r--r-- 1 root root 431M Oct 27 08:39 pytorch_model.bin
```

```

-rw-r--r-- 1 root root 518K Oct 27 08:39 vocab.txt
-rw-r--r-- 1 root root 782K Oct 27 08:39 tokenizer.json
drwxr-xr-x 2 root root 4.0K Oct 27 08:39 .
-rw-r--r-- 1 root root 862M Oct 27 08:39 training_args.bin
drwxr-xr-x 4 root root 4.0K Oct 27 09:26 ..

greekbert_model_save/2021-10-27_092610:
total 1.3G
drwxr-xr-x 4 root root 4.0K Oct 27 09:26 ..
-rw-r--r-- 1 root root 732 Oct 27 09:26 config.json
-rw-r--r-- 1 root root 431M Oct 27 09:26 pytorch_model.bin
-rw-r--r-- 1 root root 526 Oct 27 09:26 tokenizer_config.json
-rw-r--r-- 1 root root 112 Oct 27 09:26 special_tokens_map.json
-rw-r--r-- 1 root root 518K Oct 27 09:26 vocab.txt
-rw-r--r-- 1 root root 782K Oct 27 09:26 tokenizer.json
drwxr-xr-x 2 root root 4.0K Oct 27 09:26 .
-rw-r--r-- 1 root root 862M Oct 27 09:26 training_args.bin

```

▼ Part 4: Test

```

start = time.time()

epoch_loss = 0
epoch_acc = 0
predicted_labels = []
labels = []
model.eval()
for batch in test_dataloader:
    b_input_ids = batch[0].to(device)
    b_input_mask = batch[1].to(device)
    b_labels = batch[2].to(device)
    with torch.no_grad():
        loss, logits = model(
            b_input_ids,
            token_type_ids=None,
            attention_mask=b_input_mask,

```

```

        labels=b_labels,
        return_dict=False
    )
    epoch_loss += loss.item()
    logits = logits.detach().cpu().numpy()
    label_ids = b_labels.to('cpu').numpy()
    epoch_acc += flat_accuracy(logits, label_ids)
    predicted_labels.extend(np.argmax(logits, axis=1).flatten())
    labels.extend(label_ids)
    test_loss = epoch_loss / len(test_dataloader)
    test_acc = epoch_acc / len(test_dataloader)
    test_labels = labels
    test_predictions = predicted_labels


print(sklearn.metrics.classification_report(test_labels, test_predictions, target_names=["not-fake", "fake"]))
cm = sklearn.metrics.confusion_matrix(test_labels, test_predictions, normalize='true')
df_cm = pd.DataFrame(cm, index=["not-fake", "fake"], columns=["not-fake", "fake"])
seaborn.heatmap(df_cm, annot=True)
fpr, tpr, thresholds = sklearn.metrics.roc_curve(test_labels, predicted_labels)
auc = sklearn.metrics.auc(fpr, tpr)
plt.figure(figsize=(15, 5))
plt.plot(fpr, tpr, lw=3, color='darkorange', label="ROC curve (area = %.2lf)" % (auc))
plt.plot([0, 1], [0, 1], color="blue", label="Random Classifier")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.title("ROC Curve")
plt.legend()

print("Time elapsed: %.2lf" % (time.time() - start))

```

	precision	recall	f1-score	support
not-fake	0.91	0.96	0.94	107
fake	0.95	0.89	0.92	93
accuracy			0.93	200
macro avg	0.93	0.93	0.93	200
weighted avg	0.93	0.93	0.93	200

Time elapsed: 14.05



✓ 15s completed at 12:54 PM

