

## ▼ Part 1 : Load Dataset

```
from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
!pip install transformers
!pip install pytorch_wrapper
!pip install seaborn
!nvidia-smi
```

```
from google.colab import drive
drive.mount('/content/drive')
!cp /content/drive/MyDrive/fake.txt fake.txt
!cp /content/drive/MyDrive/non_fake.txt non_fake.txt
!git clone https://github.com/nlpaueb/greek-bert.git
```

```
Requirement already satisfied: filelock in /usr/local/lib/python3.7/dist-packages (from transformers) (3.3.0)
Requirement already satisfied: tokenizers<0.11,>=0.10.1 in /usr/local/lib/python3.7/dist-packages (from transformers) (0.10.1)
Requirement already satisfied: sacremoses in /usr/local/lib/python3.7/dist-packages (from transformers) (0.0.46)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from huggingface-hub<1.0.0, >=0.9.0) (4.1.1)
Requirement already satisfied: pyparsing>=2.0.2 in /usr/local/lib/python3.7/dist-packages (from packaging>=20.0->transformers) (3.0.7)
Requirement already satisfied: zipp>=0.5 in /usr/local/lib/python3.7/dist-packages (from importlib-metadata->transformers) (3.6.0)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests->transformers) (3.10.1)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests->transformers) (3.0.4)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.7/dist-packages (from requests->transformers) (2022.9.24)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.7/dist-packages (from requests->transformers) (1.26.13)
Requirement already satisfied: click in /usr/local/lib/python3.7/dist-packages (from sacremoses->transformers) (7.1.2)
Requirement already satisfied: joblib in /usr/local/lib/python3.7/dist-packages (from sacremoses->transformers) (1.1.0)
Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from sacremoses->transformers) (1.16.0)
Requirement already satisfied: pytorch_wrapper in /usr/local/lib/python3.7/dist-packages (1.1.0)
Requirement already satisfied: hyperopt<0.2,>=0.1 in /usr/local/lib/python3.7/dist-packages (from pytorch_wrapper) (0.1.2)
Requirement already satisfied: torch<2,>=1 in /usr/local/lib/python3.7/dist-packages (from pytorch_wrapper) (1.9.0)
Requirement already satisfied: scikit-learn<1,>=0.19 in /usr/local/lib/python3.7/dist-packages (from pytorch_wrapper) (0.24.2)
```

Requirement already satisfied: numpy<2,>=1.14 in /usr/local/lib/python3.7/dist-packages (from pytorch\_wrapper) (1.19.5)  
 Requirement already satisfied: tqdm<5,>=4.36.1 in /usr/local/lib/python3.7/dist-packages (from pytorch\_wrapper) (4.64.0)  
 Requirement already satisfied: six<2,>=1.11 in /usr/local/lib/python3.7/dist-packages (from pytorch\_wrapper) (1.15.0)  
 Requirement already satisfied: future in /usr/local/lib/python3.7/dist-packages (from hyperopt<0.2,>=0.1->pytorch\_wrapper) (0.18.2)  
 Requirement already satisfied: networkx in /usr/local/lib/python3.7/dist-packages (from hyperopt<0.2,>=0.1->pytorch\_wrapper) (2.6.3)  
 Requirement already satisfied: scipy in /usr/local/lib/python3.7/dist-packages (from hyperopt<0.2,>=0.1->pytorch\_wrapper) (1.7.3)  
 Requirement already satisfied: pymongo in /usr/local/lib/python3.7/dist-packages (from hyperopt<0.2,>=0.1->pytorch\_wrapper) (3.10.0)  
 Requirement already satisfied: joblib<=0.11 in /usr/local/lib/python3.7/dist-packages (from scikit-learn<1,>=0.19->pytorch\_wrapper) (1.1.0)  
 Requirement already satisfied: typing-extensions in /usr/local/lib/python3.7/dist-packages (from torch<2,>=1->pytorch\_wrapper) (4.1.1)  
 Requirement already satisfied: seaborn in /usr/local/lib/python3.7/dist-packages (0.11.2)  
 Requirement already satisfied: matplotlib<=2.2 in /usr/local/lib/python3.7/dist-packages (from seaborn) (3.2.2)  
 Requirement already satisfied: pandas<=0.23 in /usr/local/lib/python3.7/dist-packages (from seaborn) (1.1.5)  
 Requirement already satisfied: scipy<=1.0 in /usr/local/lib/python3.7/dist-packages (from seaborn) (1.4.1)  
 Requirement already satisfied: numpy<=1.15 in /usr/local/lib/python3.7/dist-packages (from seaborn) (1.19.5)  
 Requirement already satisfied: kiwisolver<=1.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib<=2.2->seaborn) (1.0.1)  
 Requirement already satisfied: pyparsing!=2.0.4,!=2.1.2,!=2.1.6,>=2.0.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib<=2.2->seaborn) (3.0.7)  
 Requirement already satisfied: cycycler<=0.10 in /usr/local/lib/python3.7/dist-packages (from matplotlib<=2.2->seaborn) (0.10.0)  
 Requirement already satisfied: python-dateutil<=2.1 in /usr/local/lib/python3.7/dist-packages (from matplotlib<=2.2->seaborn) (2.8.1)  
 Requirement already satisfied: six in /usr/local/lib/python3.7/dist-packages (from cycycler<=0.10->matplotlib<=2.2->seaborn) (1.15.0)  
 Requirement already satisfied: pytz<=2017.2 in /usr/local/lib/python3.7/dist-packages (from pandas<=0.23->seaborn) (2017.2)

Fri Nov 5 05:17:07 2021

NVIDIA-SMI 495.44				Driver Version: 460.32.03		CUDA Version: 11.2	
GPU	Name	Persistence-M	Bus-Id	Disp.A	Volatile	Uncorr. ECC	
Fan	Temp	Perf	Pwr:Usage/Cap	Memory-Usage	GPU-Util	Compute M.	MIG M.
0	Tesla K80	Off	00000000:00:04.0	Off	0		
N/A	35C	P8	29W / 149W	0MiB / 11441MiB	0%	Default	N/A

Processes:							
GPU	GI	CI	PID	Type	Process name	GPU Memory	
ID	ID	ID				Usage	
No running processes found							

Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.mount("/content/drive", force\_remount=True)

---

```
import datetime
import os
import random
import time
import unicodedata

import numpy as np
import pandas as pd
import torch
import matplotlib
import matplotlib.pyplot as plt
import transformers
from transformers import AutoTokenizer, AutoModel, get_linear_schedule_with_warmup
import torch.nn.functional as F
import sklearn
import seaborn
```

```
plt.rcParams['figure.figsize'] = [20, 5]
```

```
print("Seaborn version: %s" % (seaborn.__version__))
print("Sklearn version: %s" % (sklearn.__version__))
print("Transformers version: %s" % (transformers.__version__))
print("PyTorch version: %s" % (torch.__version__))
print("Matplotlib version: %s" % (matplotlib.__version__))
print("NumPy version: %s" % (np.__version__))
print("Pandas version: %s" % (pd.__version__))
```

```
Seaborn version: 0.11.2
Sklearn version: 0.22.2.post1
Transformers version: 4.12.3
PyTorch version: 1.9.0+cu111
Matplotlib version: 3.2.2
NumPy version: 1.19.5
Pandas version: 1.1.5
```

```
data = {'text': list(), 'label': list()}
with open('fake.txt') as infile:
    for line in infile:
        data['text'].append(line)
        data['label'].append(1)
with open('non_fake.txt') as infile:
    for line in infile:
        data['text'].append(line)
        data['label'].append(0)

df = pd.DataFrame(data)
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2201 entries, 0 to 2200
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0    text    2201 non-null    object
1    label    2201 non-null    int64
dtypes: int64(1), object(1)
memory usage: 34.5+ KB
```

```
#plot labels
labels = df.label
plt.hist(labels)
plt.xlabel("label")
plt.ylabel("number of samples")
plt.title("Labels Distribution")
plt.xticks(np.arange(len(np.unique(labels))))
```

```
([<matplotlib.axis.XTick at 0x7fa5882e0c90>,
 <matplotlib.axis.XTick at 0x7fa5865de5d0>],
 <a list of 2 Text major ticklabel objects>)
```



```
# plot article lengths and discarded data length
lengths = pd.DataFrame([len(x.split()) for x in df.text if isinstance(x, str)])
lost_data = pd.DataFrame([max(len(x.split())-512,0) for x in df.text if isinstance(x, str)])
lost_data_ratios = pd.DataFrame([max(len(x.split())-512,0)/float(len(x.split())) for x in df.text if isinstance(x, str)])

print("Tweet lengths description: %s\n" % (lengths.describe()))
lengths.plot(xlabel="tweet id", ylabel="word count", title="Tweet Length in Words", legend=None)
print("Tweet discarded data description: %s\n" % (lost_data.describe()))
lost_data.plot(xlabel="tweet id", ylabel="words removed", title="Number of Deleted Words per Tweet", legend=None)
print("Tweet discarded data ratios description: %s\n" % (lost_data_ratios.describe()))
lost_data_ratios.plot(xlabel="tweet id", ylabel="words removed ratio", title="Ratio of Deleted Words per Tweet", legend=None)

s1 = lost_data.sum()
s2 = lengths.sum()
ratio = float(s1) / float(s2)
s3 = (lengths <= 512).astype(int).sum(axis=0)
ratio2 = float(s3) / float(len(lengths))
print("Tweets exceeding 512 word limit: %d out of %d (%.2lf%%)" % (s3, len(lengths), ratio2*100))
print("Total data loss: %d out of %d words (%.2lf%% total, %.2lf%% average)" % (s1, s2, ratio*100, lost_data_ratios.mean(
```

Tweet lengths description: 0

count	2201.000000
mean	35.228987
std	21.252279
min	6.000000
25%	16.000000
50%	30.000000
75%	52.000000
max	94.000000


Tweet discarded data description: 0


count	2201.0
mean	0.0
std	0.0
min	0.0
25%	0.0
50%	0.0
75%	0.0
max	0.0

Tweet discarded data ratios description: 0


count	2201.0
mean	0.0
std	0.0
min	0.0
25%	0.0
50%	0.0
75%	0.0
max	0.0

Tweets exceeding 512 word limit: 2201 out of 2201 (100.00%)  
Total data loss: 0 out of 77539 words (0.00% total, 0.00% average)






Number of Deleted Words per Tweet



Ratio of Deleted Words per Tweet



## ▼ Part 2: Tokenization and Input Formatting

```
start = time.time()

def strip_accents_and_lowercase(s):
    return ''.join(c for c in unicodedata.normalize('NFD', s)
                    if unicodedata.category(c) != 'Mn').lower()
#####
class BinaryGreekBERTModel(torch.nn.Module):

    def __init__(self, bert_model, dp):
        super(BinaryGreekBERTModel, self).__init__()
        self._bert_model = bert_model
        self._dp = torch.nn.Dropout(dp)
        self._output_linear = torch.nn.Linear(768, 1)

    def forward(self, text, text_len):
        attention_mask = pwF.create_mask_from_length(text_len, text.shape[1])
        return self._output_linear(self._dp(self._bert_model(text, attention_mask=attention_mask)[0][:, 0, :]))
#####

#from transformers import BertTokenizer
#tokenizer = BertTokenizer.from_pretrained('bert-base-uncased', do_lower_case=True)
tokenizer = AutoTokenizer.from_pretrained("nlpueb/bert-base-greek-uncased-v1")
greekbert_model = AutoModel.from_pretrained('nlpueb/bert-base-greek-uncased-v1')
binary_greekbert_model = BinaryGreekBERTModel(greekbert_model, dp=0)

working_set_df = df.head(2200) # 2201 items don't split nicely, discard one

text_data = working_set_df.text
```



```

text_data.fillna('', inplace=True)

input_ids = []
attention_masks = []
for item in text_data:
    unaccented_string = strip_accents_and_lowercase(item)
    encoded_dict = tokenizer.encode_plus(unaccented_string,
                                         add_special_tokens=True,
                                         max_length=100,
                                         truncation=True,
                                         padding='max_length',
                                         #pad_to_max_length=True,
                                         return_attention_mask=True,
                                         return_tensors="pt")

    input_ids.append(encoded_dict["input_ids"])
    attention_masks.append(encoded_dict["attention_mask"])

input_ids = torch.cat(input_ids, dim=0)
attention_masks = torch.cat(attention_masks, dim=0)
labels = torch.tensor(working_set_df.label.values)

print("Time elapsed: %.1lf seconds." % (time.time() - start))

```

Some weights of the model checkpoint at nlpaueb/bert-base-greek-uncased-v1 were not used when initializing BertModel  
- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with  
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactl  
/usr/local/lib/python3.7/dist-packages/pandas/core/series.py:4536: SettingWithCopyWarning:  
A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: [https://pandas.pydata.org/pandas-docs/stable/user\\_guide/indexing.html#returnir](https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returnir)  
downcast=downcast,  
Time elapsed: 4.7 seconds.

```

# split training and validation sets
from torch.utils.data import TensorDataset, random_split, DataLoader, RandomSampler, SequentialSampler

dataset = TensorDataset(input_ids, attention_masks, labels)
train_size = int(0.8 * len(dataset))

```

```

train_size = int(0.8 * len(dataset))
validation_size = int(0.1 * len(dataset))
test_size = int(0.1 * len(dataset))
train_dataset, validation_dataset, test_dataset = random_split(dataset, [train_size, validation_size, test_size])
print("Training size: %s, Validation size: %s Test size: %s" % (train_size, validation_size, test_size))

batch_size = 8
train_dataloader = DataLoader(train_dataset, sampler=RandomSampler(train_dataset), batch_size=batch_size)
validation_dataloader = DataLoader(validation_dataset, sampler=SequentialSampler(validation_dataset), batch_size=batch_size)
test_dataloader = DataLoader(test_dataset, sampler=SequentialSampler(test_dataset), batch_size=batch_size)

```

Training size: 1760, Validation size: 220 Test size: 220

## ▼ Part 3: Training

```

start = time.time()

from transformers import BertForSequenceClassification, AdamW, BertConfig

model = BertForSequenceClassification.from_pretrained(
    "nlpau/bert-base-greek-uncased-v1",
    num_labels = 2,
    output_attentions = False,
    output_hidden_states = False,
)

device = None
if torch.cuda.is_available():
    device = torch.device("cuda:0")
    model.cuda()
else:
    device = torch.device("cpu")

#####
# Batch sizes: 16, 32

```

```
# Learning rate (Adam): 5e-5, 3e-5, 2e-5
# Epochs: 2, 3, 4
#####

optimizer = AdamW(model.parameters(), lr=2e-5, eps=1e-8)
epochs = 4
total_steps = len(train_dataloader) * epochs
scheduler = get_linear_schedule_with_warmup(optimizer,
                                             num_warmup_steps = 0, # Default value in run_glue.py
                                             num_training_steps = total_steps)

print("Time elapsed: %.2lf" % (time.time() - start))
!nvidia-smi
```

Some weights of the model checkpoint at `nlpaueb/bert-base-greek-uncased-v1` were not used when initializing BertForSe

- This IS expected if you are initializing BertForSequenceClassification from the checkpoint of a model trained on a
- This IS NOT expected if you are initializing BertForSequenceClassification from the checkpoint of a model that you

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at `nlpaueb/bert-base-gr`

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

Time elapsed: 21.27

Fri Nov 5 05:24:17 2021

+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
NVIDIA-SMI 495.44			Driver Version: 460.32.03				CUDA Version: 11.2		
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									
GPU	Name		Persistence-M	Bus-Id	Disp.A	Volatile Uncorr. ECC			
Fan	Temp	Perf	Pwr:Usage/Cap		Memory-Usage	GPU-Util	Compute M.	MIG M.	
=====+=====+=====+=====+=====+=====+=====+=====+=====+=====									
0	Tesla K80		Off	00000000:00:04.0	Off			0	
N/A	40C	P0	60W / 149W	1039MiB / 11441MiB		27%	Default	N/A	
+-----+-----+-----+-----+-----+-----+-----+-----+-----+-----+									

```
+-----+
| Processes:                                     |
| GPU    GI    CI          PID   Type   Process name                      GPU Memory |
|      ID    ID                          Usage                      |
|=====|
| No running processes found                    |
+-----+
```

## ▼ Training Loop

```
start = time.time()

def flat_accuracy(preds, labels):
    """Function to calculate the accuracy of our predictions vs labels"""
    pred_flat = np.argmax(preds, axis=1).flatten()
    labels_flat = labels.flatten()
    return np.sum(pred_flat == labels_flat) / len(labels_flat)
#####

seed_val = 21
random.seed(seed_val)
np.random.seed(seed_val)
torch.manual_seed(seed_val)
torch.cuda.manual_seed_all(seed_val)

training_stats = []
total_t0 = time.time()

for epoch_i in range(0, epochs):
    print('Epoch %s out of %s' % (epoch_i + 1, epochs))
    print('Training...')

    t0 = time.time()
    total_train_loss = 0
    model.train()
    for step, batch in enumerate(train_dataloader):
        if step % 40 == 0 and not step == 0:
            print(' Batch %s out of %s, Elapsed %.2lf seconds' % (step, len(train_dataloader), time.time() - t0))

        b_input_ids = batch[0].to(device)
        b_input_mask = batch[1].to(device)
        b_labels = batch[2].to(device)
```

```

model.zero_grad()
loss, logits = model(
    b_input_ids,
    token_type_ids=None,
    attention_mask=b_input_mask,
    labels=b_labels,
    return_dict=False
)

total_train_loss += loss.item()
loss.backward()
torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)
optimizer.step()
scheduler.step()
avg_train_loss = total_train_loss / len(train_dataloader)
training_time = time.time() - t0

print(" Average training loss: %.2lf" % (avg_train_loss))
print(" Epoch %s training time %.2lf seconds" % (epoch_i, training_time))
print("Running Validation...")

t0 = time.time()
model.eval()
total_eval_accuracy = 0
total_eval_loss = 0
nb_eval_steps = 0
for batch in validation_dataloader:
    b_input_ids = batch[0].to(device)
    b_input_mask = batch[1].to(device)
    b_labels = batch[2].to(device)
    with torch.no_grad():
        loss, logits = model(
            b_input_ids,
            token_type_ids=None,
            attention_mask=b_input_mask,
            labels=b_labels,
            return_dict=False
        )

```

```

        total_eval_loss += loss.item()
        logits = logits.detach().cpu().numpy()
        label_ids = b_labels.to('cpu').numpy()
        total_eval_accuracy += flat_accuracy(logits, label_ids)
    avg_val_accuracy = total_eval_accuracy / len(validation_dataloader)
    avg_val_loss = total_eval_loss / len(validation_dataloader)
    validation_time = time.time() - t0

    print(" Accuracy: %.2lf" % (avg_val_accuracy))
    print(" Validation Loss: %2lf" % (avg_val_loss))
    print(" Validation time %.2lf seconds" % (validation_time))

    training_stats.append(
        {
            'epoch': epoch_i + 1,
            'Training Loss': avg_train_loss,
            'Valid. Loss': avg_val_loss,
            'Valid. Accur.': avg_val_accuracy,
            'Training Time': training_time,
            'Validation Time': validation_time
        }
    )
print("Training complete!")
print("Time elapsed: %.2lf seconds" % (time.time() - start))

```

Epoch 1 out of 4

Training...

```

Batch 40 out of 220, Elapsed 14.14 seconds
Batch 80 out of 220, Elapsed 27.84 seconds
Batch 120 out of 220, Elapsed 41.57 seconds
Batch 160 out of 220, Elapsed 55.52 seconds
Batch 200 out of 220, Elapsed 69.50 seconds
Average training loss: 0.47
Epoch 0 training time 76.36 seconds

```

Running Validation...

```

Accuracy: 0.83
Validation Loss: 0.352907

```

Validation time 3.01 seconds  
Epoch 2 out of 4  
Training...  
Batch 40 out of 220, Elapsed 13.78 seconds  
Batch 80 out of 220, Elapsed 27.54 seconds  
Batch 120 out of 220, Elapsed 41.27 seconds  
Batch 160 out of 220, Elapsed 55.28 seconds  
Batch 200 out of 220, Elapsed 69.09 seconds  
Average training loss: 0.29  
Epoch 1 training time 75.82 seconds  
Running Validation...  
Accuracy: 0.90  
Validation Loss: 0.278729  
Validation time 2.95 seconds  
Epoch 3 out of 4  
Training...  
Batch 40 out of 220, Elapsed 13.51 seconds  
Batch 80 out of 220, Elapsed 27.03 seconds  
Batch 120 out of 220, Elapsed 40.48 seconds  
Batch 160 out of 220, Elapsed 53.94 seconds  
Batch 200 out of 220, Elapsed 67.49 seconds  
Average training loss: 0.21  
Epoch 2 training time 74.28 seconds  
Running Validation...  
Accuracy: 0.93  
Validation Loss: 0.308774  
Validation time 2.97 seconds  
Epoch 4 out of 4  
Training...  
Batch 40 out of 220, Elapsed 13.51 seconds  
Batch 80 out of 220, Elapsed 26.92 seconds  
Batch 120 out of 220, Elapsed 40.33 seconds  
Batch 160 out of 220, Elapsed 53.72 seconds  
Batch 200 out of 220, Elapsed 67.19 seconds  
Average training loss: 0.12  
Epoch 3 training time 73.88 seconds  
Running Validation...  
Accuracy: 0.92  
Validation Loss: 0.377954  
Validation time 2.98 seconds  
Training complete!  
Time elapsed: 312.27 seconds

!nvidia-smi

Fri Nov 5 05:30:50 2021

+-----+-----+-----+-----+-----+-----+									
NVIDIA-SMI 495.44			Driver Version: 460.32.03				CUDA Version: 11.2		
+-----+-----+-----+-----+-----+-----+									
GPU	Name		Persistence-M		Bus-Id	Disp.A	Volatile	Uncorr.	ECC
Fan	Temp	Perf	Pwr:Usage/Cap			Memory-Usage	GPU-Util	Compute M.	M.
=====									
0	Tesla K80		Off		00000000:00:04.0	Off			0
N/A	57C	P0	61W / 149W		3230MiB / 11441MiB		0%	Default	N/A
+-----+-----+-----+-----+-----+-----+									

+-----+-----+-----+-----+-----+-----+									
Processes:									
GPU	GI	CI	PID	Type	Process name	GPU Memory			
	ID	ID				Usage			
=====									
No running processes found									
+-----+-----+-----+-----+-----+-----+									

```
#####  
# View training summary  
#####
```

```
pd.set_option('precision', 2)  
df_stats = pd.DataFrame(data=training_stats)  
df_stats = df_stats.set_index('epoch')  
df_stats
```



	Training Loss	Valid. Loss	Valid. Accur.	Training Time	Validation Time
epoch					
1	0.47	0.35	0.83	76.36	3.01

```
# If validation loss is increasing, it means we are overfitting!
# Plot training and validation loss
import seaborn as sns
```

```
# Use plot styling from seaborn.
sns.set(style='darkgrid')
```

```
# Increase the plot size and font size.
sns.set(font_scale=1.5)
plt.rcParams["figure.figsize"] = (12,6)
```

```
# Plot the learning curve.
plt.plot(df_stats['Training Loss'], 'b-o', label="Training")
plt.plot(df_stats['Valid. Loss'], 'g-o', label="Validation")
```

```
# Label the plot.
plt.title("Training & Validation Loss")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.xticks([1, 2, 3, 4])
```

```
plt.show()
```



```
#####
```

```
# Save model to disk
```

```
#####
```

```
tstamp = datetime.datetime.now().strftime("%Y-%m-%d_%H%M%S")
```

```
output_dir = './greekbert_twitter_model_save/%s/' % (tstamp)
```

```
if not os.path.exists(output_dir):
```

```
    os.makedirs(output_dir)
```

```
print("Saving model to %s" % output_dir)
```

```
model_to_save = model.module if hasattr(model, 'module') else model # Take care of distributed/parallel training
```

```
model_to_save.save_pretrained(output_dir)
```

```
tokenizer.save_pretrained(output_dir)
```

```
torch.save(optimizer.state_dict(), os.path.join(output_dir, 'training_args.bin'))
```

```
Saving model to ./greekbert_twitter_model_save/2021-11-05_053152/
```

```
!ls -lhatrR greekbert_twitter_model_save/
```

```
greekbert_twitter_model_save/:
```

```
total 12K
```

```
drwxr-xr-x 1 root root 4.0K Nov  5 05:31 ..
```

```
drwxr-xr-x 3 root root 4.0K Nov  5 05:31 .
```

```
drwxr-xr-x 2 root root 4.0K Nov  5 05:31 2021-11-05_053152
```

```

greekbert_twitter_model_save/2021-11-05_053152:
total 1.3G
drwxr-xr-x 3 root root 4.0K Nov  5 05:31 ..
-rw-r--r-- 1 root root  732 Nov  5 05:31 config.json
-rw-r--r-- 1 root root 431M Nov  5 05:31 pytorch_model.bin
-rw-r--r-- 1 root root  526 Nov  5 05:31 tokenizer_config.json
-rw-r--r-- 1 root root  112 Nov  5 05:31 special_tokens_map.json
-rw-r--r-- 1 root root 518K Nov  5 05:31 vocab.txt
-rw-r--r-- 1 root root 782K Nov  5 05:31 tokenizer.json
drwxr-xr-x 2 root root 4.0K Nov  5 05:31 .
-rw-r--r-- 1 root root 862M Nov  5 05:31 training_args.bin

```

## ▼ Part 4: Test

```

start = time.time()

epoch_loss = 0
epoch_acc = 0
predicted_labels = []
labels = []
model.eval()
for batch in test_dataloader:
    b_input_ids = batch[0].to(device)
    b_input_mask = batch[1].to(device)
    b_labels = batch[2].to(device)
    with torch.no_grad():
        loss, logits = model(
            b_input_ids,
            token_type_ids=None,
            attention_mask=b_input_mask,
            labels=b_labels,
            return_dict=False
        )
    epoch_loss += loss.item()
    logits = logits.detach().cpu().numpy()

```



```
label_ids = b_labels.to('cpu').numpy()
epoch_acc += flat_accuracy(logits, label_ids)
predicted_labels.extend(np.argmax(logits, axis=1).flatten())
labels.extend(label_ids)
test_loss = epoch_loss / len(test_dataloader)
test_acc = epoch_acc / len(test_dataloader)
test_labels = labels
test_predictions = predicted_labels


print(sklearn.metrics.classification_report(test_labels, test_predictions, target_names=["not-fake", "fake"]))
cm = sklearn.metrics.confusion_matrix(test_labels, test_predictions, normalize='true')
df_cm = pd.DataFrame(cm, index=["not-fake", "fake"], columns=["not-fake", "fake"])
seaborn.heatmap(df_cm, annot=True)
fpr, tpr, thresholds = sklearn.metrics.roc_curve(test_labels, predicted_labels)
auc = sklearn.metrics.auc(fpr, tpr)
plt.figure(figsize=(15, 5))
plt.plot(fpr, tpr, lw=3, color='darkorange', label="ROC curve (area = %.2lf)" % (auc))
plt.plot([0, 1], [0, 1], color="blue", label="Random Classifier")
plt.xlabel("False positive rate")
plt.ylabel("True positive rate")
plt.title("ROC Curve")
plt.legend()

print("Time elapsed: %.2lf" % (time.time() - start))
```

	precision	recall	f1-score	support
not-fake	0.93	0.94	0.94	148
fake	0.87	0.86	0.87	72
accuracy			0.91	220
macro avg	0.90	0.90	0.90	220
weighted avg	0.91	0.91	0.91	220

Time elapsed: 3.29





✓ 4s completed at 7:32 AM

